

Article

Beyond Black-Box Deep Knowledge Tracing: Transformers with Representational Grounding for Pedagogical Interpretability

Concha Labra ¹, Olga C. Santos ^{2,*}

¹ Department of Artificial Intelligence, Computer Science School, UNED, 28040 Madrid, Spain; clabra@dia.uned.es

² PhyUM Research Center, Department of Artificial Intelligence, Computer Science School, UNED, 28040 Madrid, Spain; ocsantos@dia.uned.es

* Author to whom correspondence should be addressed.

Abstract

This study introduces iDKT, an interpretable-by-design Transformer model that utilizes *Representational Grounding* to align deep latent representations with educational constructs, leveraging the high accuracy of deep knowledge tracing models while addressing their inherent lack of interpretability. We introduce a formal validation framework to verify the alignment of iDKT's internal representations and, using Bayesian Knowledge Tracing (BKT) as a reference, evaluate the model across multiple educational datasets. Results demonstrate that iDKT maintains state-of-the-art predictive performance while yielding additional interpretable insights at a significantly higher granularity than those provided by the reference model. Specifically, iDKT identifies student-level initial knowledge and learning velocities, providing mastery estimates that are more sensitive to the nuances of individual behavioral patterns than those produced by standard BKT. These individualized insights enable precise diagnostic placement and dynamic pacing, allowing adaptive learning environments to tailor instruction to each student's unique learning profile with enhanced precision. This work offers both a robust methodology for evaluating the interpretability of Transformer-based models and a practical tool for improving educational effectiveness through data-driven personalization.

Keywords: deep knowledge tracing; transformer; interpretability; Bayesian Knowledge Tracing; educational data analysis; personalized learning

1. Introduction

Knowledge Tracing [1] is a fundamental task in the fields of Artificial Intelligence in Education, Intelligent Tutoring Systems and Massive Open Online Courses. Its primary objective is to model a student's dynamic knowledge state over time based on their history of interactions with learning materials, enabling systems to predict future performance and provide personalized instruction. As educational environments become increasingly diverse and digital, the ability to accurately track and interpret student mastery has become a critical requirement for scalable, effective education.

Historically, the field has been dominated by two distinct paradigms. The first, exemplified by Bayesian Knowledge Tracing (BKT) and its variants [2], relies on probabilistic graphical models that explicitly represent knowledge states. BKT models are intrinsically interpretable, being based on parameters such as initial knowledge, learning rate, or slipping and guessing probabilities that map directly to pedagogical constructs, allowing educators

Received:

Revised:

Accepted:

Published:

Copyright: © 2025 by the authors.

Submitted to *Journal Not Specified* for possible open access publication under the terms and conditions of the Creative Commons Attribution (CC BY) license.

to understand how they work and trust their decisions. However, this interpretability comes at the cost of a simplicity that often limits its predictive power, making them struggle to capture the complex, non-linear dependencies often present in educational datasets.

The second paradigm emerged with the advent of Deep Knowledge Tracing (DKT) [3], which uses different variants of deep learning techniques from the initial Recurrent Neural Networks to current Transformers [4] to model student interactions. These models have achieved state-of-the-art predictive performance, significantly outperforming classical approaches by leveraging the high capacity of deep learning models that allows them to learn complex patterns [5]. Yet, this predictive power has come at a significant cost: interpretability. Deep learning models are notoriously opaque "black boxes," where the learned representations are distributed across high-dimensional latent spaces that bear no direct correspondence to constructs with a clear semantic meaning. This lack of transparency creates a trust gap for practitioners, who cannot easily discern why a model predicts a student has failed or succeeded, nor can derive actionable pedagogical insights from the model's internal state [6].

Current efforts to bridge this gap typically rely on post-hoc explainability methods, such as weights visualization or perturbation analysis [7,8]. While valuable for debugging, these techniques often provide only a superficial view of the model's decision-making process and do not guarantee that the learned representations align with meaningful constructs. Moreover, their application and interpretation require technical deep learning expertise, limiting their accessibility to practitioners without this specialized knowledge.

To address these limitations, we propose a shift towards interpretability-by-design, inspired by the emerging paradigm of Theory-Guided Data Science (TGDS) [9]. In TGDS, maintaining consistency with theoretical postulates is an architectural constraint rather than an afterthought. By integrating extensive domain knowledge, TGDS-based models can be constrained to learn representations that are both theoretically plausible and highly predictive. While this approach has been applied mostly to science—and specifically to physics [10]—we adapt it here to the educational domain.

Standard TGDS implementations typically rely on auxiliary loss functions to incorporate formal knowledge expressed as rules, algebraic constraints, or differential equations [11]. We propose a novel approach called *Representational Grounding* that, in contrast, utilizes auxiliary losses operating on projections of the Transformer's embeddings. This mechanism enables the model to learn representations that are consistent with semantically meaningful constructs.

The major contributions of this work are as follows:

- Proposal of Representational Grounding, a novel method that overcomes the black-box nature of Transformers by providing interpretability-by-design.
- Introduction of a formal validation framework to quantify interpretability via representational alignment, enabling a systematic characterization of the trade-off between reference fidelity and predictive performance.
- Application of Representational Grounding to the development of iDKT, a new type of knowledge tracing models that leverage the high accuracy inherent in deep learning while achieving pedagogical interpretability.
- Empirical demonstration of iDKT benefits by showing how it captures granular, student-specific insights—such as individualized initial knowledge and learning rates—that are beyond the capabilities of simpler models such as BKT.

The remainder of this paper is structured as follows. Section 2 reviews related work on knowledge tracing, deep learning interpretability, and theory-guided data science. Section 3 describes the proposed iDKT architecture and the *Representational Grounding* framework. Sections 4 and 5 detail the individualized embedding mechanism and the multi-objective

loss system, respectively. Section 6 presents the results of the experimental validation, interpretability results, and the analysis of individualization granularity. Finally, Section 7 concludes the paper.

2. Related Work

- 2.1. Deep Knowledge Tracing
- 2.2. Bayesian Knowledge Tracing
- 2.2.1. Parameters of the Standard Model
- 2.2.2. Individualization
- 2.3. Deep Learning Interpretability
- 2.4. Theory-Guided Data Science

3. iDKT Model

We propose an Interpretable Deep Knowledge Tracing (iDKT) model with a Transformer-based architecture designed to bridge the gap between the high predictive capacity of deep learning and the intrinsic interpretability of simpler models such as Bayesian Knowledge Tracing (BKT). Unlike standard black-box models, iDKT utilizes a novel mechanism called Representational Grounding to anchor its latent representations to the conceptual space of a interpretable model choosen as reference.

The core architecture of iDKT, as illustrated in Figure 1, extends the standard Transformer framework [4] by incorporating specialized components for Representational Grounding, primarily integrated within the embedding layers and the multi-objective loss pipeline.

1. **Input Data:** This stage handles the ingestion of student interactions (concepts, questions, and binary responses $r \in \{0, 1\}$). Crucially, it also loads values from the BKT reference model, including performance predictions (p_{BKT}) and per-skill parameters such as initial mastery (l_0) and learning transition rates (T), which serve as grounding targets for the model.
2. **Embeddings:** Task and history information are projected into a continuous embedding space. We then apply the individualization transformations described in Section 4 where base embeddings are augmented by student-specific parameters: a Proficiency Offset (lc) and a Rate Augmentation (tc). This produces the Transition Gap (individualized task representation x') and the Transition Gain (individualized interaction history y'), effectively defining the difficulty of a task relative to the student's baseline proficiency.
3. **Transformer Core:** The model employs a dual-encoder architecture followed by a cross-attention decoder. One encoder processing the individualized interaction history ($y'_{1:t-1}$) to capture global student behavior, while a task encoder processes the current individualized task (x'_t). The decoder uses multi-head cross-attention to synthesize these streams into a latent Proficiency Context, representing the student's current specialized knowledge state for the target task.
4. **Output Stage:** The latent Proficiency Context is passed through a multi-layer perceptron (MLP) that maps the deep representations to the final output space. The final output is the iDKT prediction (p_{iDKT}), representing the probability that the student will respond correctly to the current task.
5. **Loss Functions (Grounding Pipeline):** During training, the architecture is supervised not only by the prediction loss (L_{sup}) against ground truth outcomes but also by an alignment pipeline. This includes the Alignment Loss (L_{ref}) which penalizes deviations from the BKT prediction, and Parameter Losses (L_{init}, L_{rate}) that ground

the model's internal proficiency and rate parameters to their BKT-derived theoretical counterparts.

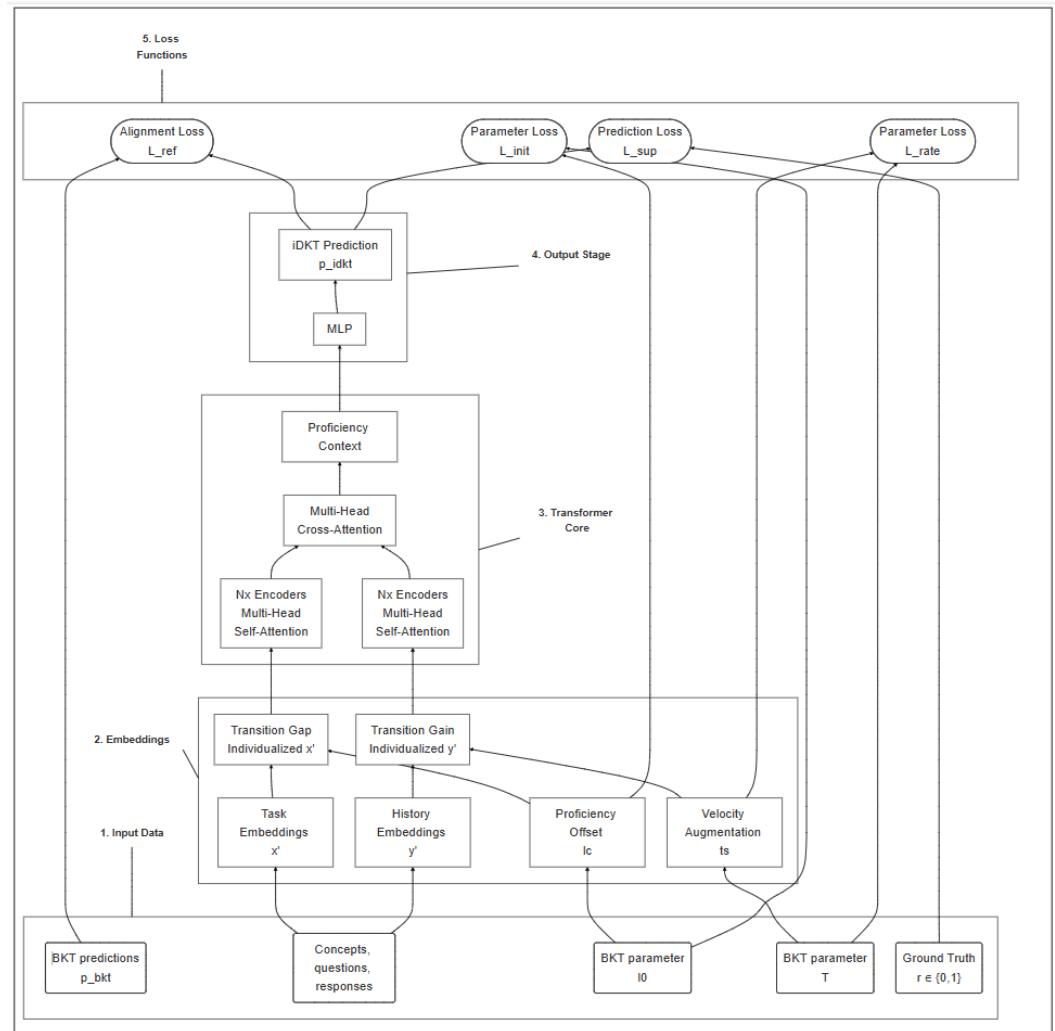
130
131

Figure 1. The iDKT Architecture. The diagram illustrates the five functional stages: (1) Input Data ingestion including BKT targets, (2) Individualized Embeddings, (3) the Transformer Core, (4) the MLP-based Output, and (5) the Loss Functions for representational grounding.

4. Embeddings

132
133
134
135
136
137
138
139
140
141
142
143
144

In standard educational datasets, such as ASSISTments 2009, ASSISTments 2015, Algebra 2005, and others [12], student interactions are recorded at the level of specific questions or tasks, each of which is associated with one or more underlying concepts or knowledge components. This structure reflects the fact that proficiency in a concept (e.g., the Pythagorean Theorem) is acquired through interactions with a diverse range of tasks. While all tasks involving a concept share the same semantic core, they differ in their specific manifestations—most notably in their intrinsic difficulty or complexity. Therefore, an effective representation must capture both the shared identity of the concept and the unique deviation of the specific task.

In Transformer-based models [4] we can operationalize this principle representing the tasks as embedding vectors:

$$x' = c + u \cdot d \quad (\text{Task}) \quad (1)$$

In this formulation, c acts as the *Concept Anchor*, a vector representing the invariant semantic identity of the concept while the vector d represents the learnable *Question Variation Axis*, defining the specific direction of the "transition gap". The scalar u serves as the *Relational Magnitude*, representing the question's specific relative difficulty compared to other questions involving the same concept. Consequently, rather than encode arbitrary embeddings for every question, we encode the vector sum of two distinct components with clear semantic meaning: a base *concept identity* and a *difficulty shift*. 145
146
147
148
149
150
151

In a similar way, we can operationalize the interactions between questions and students as embedding vectors: 152
153

$$y' = e + u \cdot (f + d) \quad (\text{Interaction History}) \quad (2) \quad 154$$

Here e represents the *Interaction Base*, which is a combined representation of concept c and the binary outcome r (correct/incorrect), while f represents the *Interaction Variation Axis*, which is similar to the Question Variation Axis (d_c) but is specific to the interaction between a question and a student. The inclusion of d_c in the interaction shift ensures that the difficulty vectors are consistent across both questions (x') and interactions (y'). 155
156
157
158
159

Extending this rationale, we can enrich the x' embeddings by integrating additional components with explicit semantic significance. Specifically, by adopting Bayesian Knowledge Tracing (BKT) as a reference model, we can incorporate vectors corresponding to its core theoretical parameters—Initial Knowledge (L_0) and Learning Rate (T)—thereby grounding the deep representation in established pedagogical constructs. 160
161
162
163
164

To get individualized values for these parameters, we decompose them into population-level bases and student-specific deviations: 165
166

$$l_c = L_0 + k_s \cdot d_k \quad (\text{Personalized Initial Knowledge}) \quad (3) \quad 167$$

$$t_c = T + v_s \cdot d_v \quad (\text{Personalized Learning Rate}) \quad (4) \quad 168 \quad 169$$

where l_c is the personalized initial knowledge for concept, t_c is the personalized learning rate for concept, L_0 and T are the population-level base embeddings, d_k and d_v are the learnable variation axes vectors (similar to the difficulty axis d), and k_s and v_s are the scalar student-specific deviations learned for each individual. 170
171
172
173

We include these vectors to get the final input embedding for the encoder and decoder components of the Transformer: 174
175

$$x' = (c + u \cdot d) - l_c \quad (\text{Individualized Task}) \quad (5) \quad 176$$

$$y' = (e + u \cdot (f + d)) + t_c \quad (\text{Individualized Interaction History}) \quad (6) \quad 177 \quad 178$$

where c represents the concept embedding, u the question-specific difficulty shift, d the task variation axis, e the interaction base, f the interaction variation axis, l_c the personalized initial knowledge, and t_c the personalized learning rate. 179
180
181

The rationale for using difference for the individualized task (x') and sum for the interaction history (y') is due to their distinct semantic roles: 182
183

- x' represents the *Transition Gap: Difficulty – Proficiency*. Under this relational logic, objective task difficulty is offset by prior proficiency, ensuring that task demands are defined relative to the subject's baseline. This formulation captures the difference between the task requirements and the current state, representing the residual gap after accounting for latent proficiency. 184
185
186
187
188

- y' represents the *Transition Gain: Interaction + Rate*. The history encoder accumulates evidence from interactions, applying a consistent relational logic where the signal value is augmented by latent rate, ensuring that interaction outcomes are defined relative to the subject's pace. Under this formulation, the total value encompasses not only the interaction outcome but also the rate of progress through the state trajectory, as this incremental gain serves as a robust indicator of future performance. This approach, therefore, captures latent progression by augmenting observed outcomes with transition rate, thereby reflecting individualized acquisition rates.

5. Loss Functions

The model is trained using a multi-objective loss function designed to ensure that the high-capacity Transformer remains aligned with pedagogical principles through Representational Grounding. The total loss \mathcal{L}_{total} is defined as a weighted sum of different loss components described in detail below.

$$\mathcal{L}_{total} = L_{sup} + \lambda_{ref} L_{ref} + \lambda_{init} L_{init} + \lambda_{rate} L_{rate} + L_{reg} \quad (7)$$

5.1. Supervised Alignment (L_{sup})

The primary objective L_{sup} uses standard Binary Cross-Entropy (BCE) between the iDKT performance predictions \hat{y}_t and the observed ground truth outcomes $r_t \in \{0, 1\}$. This loss ensures predictive accuracy by minimizing the deviance from observed student behavior:

$$L_{sup} = -\frac{1}{N} \sum_{t=1}^N [r_t \log(\hat{y}_t) + (1 - r_t) \log(1 - \hat{y}_t)] \quad (8)$$

5.2. Representational Grounding

The grounding losses ($L_{ref}, L_{init}, L_{rate}$) use Mean Squared Error (MSE) to anchor deep representations to BKT-derived values. Specifically, L_{ref} forces behavioral predictions to stay close to the theoretical baseline, while L_{init} and L_{rate} ground the individualized parameters l_c and t_c in meaningful educational starting points and acquisition paces, respectively. Instead of arbitrary latent weights, the model's internal states are projected through a sigmoid activation $\sigma(\cdot)$ and compared directly to the reference values:

$$L_{ref} = \text{MSE}(\hat{y}, p_{BKT}) \quad (9)$$

$$L_{init} = \text{MSE}(\sigma(\bar{l}_c), L0_{BKT}) \quad (10)$$

$$L_{rate} = \text{MSE}(\sigma(\bar{t}_c), T_{BKT}) \quad (11)$$

where \bar{l}_c and \bar{t}_c are the average across the feature dimension of the individualized embeddings for proficiency and rate, respectively. This formulation forces the deep representation to be not only predictive but also semantically consistent with the reference constructs.

5.3. Inductive Bias Regularization (L_{reg})

While the grounding losses anchor the global position of the latent space to the BKT parameter estimations, L_{reg} ensures that student-level individualization is *parsimonious*. This loss acts directly on the individualization parameters (u_q, k_s, v_s) to ensure that the model only deviates from the theoretical prior when functionally necessary. We apply distinct L_2 penalties to the scalar parameters governing variation:

$$L_{reg} = \lambda_u \sum_{q \in Q} u_q^2 + \lambda_k \sum_{s \in S} k_s^2 + \lambda_v \sum_{s \in S} v_s^2 \quad (12)$$

where u_q represents item difficulty, k_s is the student-specific knowledge gap, and v_s is the learning rate deviation. This formulation implements a *normal student prior*: the model assumes every subject adheres to the population-level parameters derived from the BKT reference unless their unique interaction history provides sufficient signal to justify the regularization cost.

6. Results and Discussion

6.1. Research Questions

The experimental validation of iDKT is guided by the following research questions:

1. **Interpretability Validation (RQ1):** Is it possible to rigorously validate that a iDKT model, whose representations are grounded in a reference model, actually yields interpretable constructs?
2. **Trade-Off between Predictive Performance and Interpretability (RQ2):** To what extent can deep knowledge tracing models be constrained for interpretability alignment without significantly degrading predictive performance?
3. **Improvement of Pedagogical Diagnostics (RQ3):** How can the Transformer's ability to capture longitudinal context be utilized to improve or complement pedagogical diagnostics?

6.2. Experimental Setup

We implemented the iDKT model in PyTorch and used the benchmark library PYKT [12] to leverage standardized data preprocessing, dataset splitting and benchmarking of baseline models. For training, we used standard 5-fold cross-validation with an 80/20 train/test split. The model was trained using the Adam optimizer with a learning rate of $1e - 4$, a batch size of 64, and a dropout rate of 0.2 to prevent overfitting. The maximum number of epochs was set to 200, with an early stopping mechanism (patience=10) to terminate training if validation performance plateaued.

The Transformer architecture was configured with an embedding dimension d_{model} of 256, 8 attention heads, and 4 encoder/decoder blocks. The feed-forward dimension d_{ff} was set to 512. Regularization penalties for individualization parameters (L_2 on u_q , k_s , v_s) were all set to $1e - 5$. For reproducibility, all experiments were seeded (seed=42) and executed on NVIDIA A100 GPU infrastructure. Predictive performance was evaluated using Area Under the ROC Curve (AUC) and Accuracy (ACC), while interpretability was assessed using the metrics defined in Section 6.5.

6.3. Datasets

We did the evaluation of iDKT with these 5 widely used datasets:

- ASSISTments2009: A dataset consisting of math exercises, collected from the free online tutoring ASSISTments platform in 2009-2010. It is one of the most widely used and has been the standard benchmark for many years [13].
- ASSISTments2015: This dataset was collected from the ASSISTments platform in the year of 2015. It has the largest number of students among the other ASSISTments datasets [13].
- Algebra2005: A dataset from the KDD Cup 2010 EDM Challenge containing questions from the Carnegie Learning Algebra system deployed 2005-2006 [14].
- Bridge2006: A dataset from the KDD Cup 2010 EDM Challenge with the Carnegie Learning Bridge to Algebra system, deployed 2006-2007 [15].
- NIPS34: A dataset collected from the Eedi platform [16] containing answers to multiple-choice diagnostic math questions for the Tasks 3 & 4 at the NeurIPS 2020 Education Challenge.

6.4. Predictive Performance

To evaluate the predictive performance of the iDKT model, we used the Area Under the Curve (AUC) and the Classification Accuracy (ACC) of the models on the test set with the 5 datasets described in Section 6.3. The results are shown in the Table 1.

Table 1. Predictive Performance of the iDKT model across 5 datasets.

Model	AS2009_S	AS2015	Algebra2005	Bridge2006	NIPS34
iDKT	0.8255	0.7252	0.9281	0.8092	0.7987

When comparing these results with state-of-the-art DKT models reported in [12], we observe that iDKT achieves the best predictive performance on the AS2009 and Algebra2005 datasets and ranks second on the Bridge2006 and NIPS34 datasets, surpassed only by the AKT model.

6.5. Interpretability Validation (RQ1)

To verify that iDKT's internal representations—specifically Personalized Initial Knowledge (\mathbf{l}_c) and Personalized Learning Rate (\mathbf{t}_c) as defined in Equations 3 and 4—faithfully represent the educational constructs postulated by the reference model, we employ two metrics widely utilized in psychometrics and educational measurement [17,18]:

1. *Convergent Validity (Latent Fidelity):* This metric, denoted as I_1 , is calculated as the Pearson correlation (r) between the projected latent factors ($\mathbf{l}_u, \mathbf{t}_u$) and the reference BKT parameters (L_0 and T). High alignment proves the model has successfully internalized the theoretical constructs.

The metrics are expressed as:

$$I_{1l} = \text{Corr}(l_u, L_0) \quad \text{and} \quad I_{1t} = \text{Corr}(t_u, T) \quad (13)$$

where the projected latent factors are calculated via:

$$l_u = \sigma(\bar{l}_c), \quad t_u = \sigma(\bar{t}_c) \quad (14)$$

with \bar{l}_c and \bar{t}_c representing the mean across the feature dimension of the individualized embeddings.

2. *Predictor Equivalence (Behavioral Alignment):* This metric, denoted as I_2 , assesses the functional substitutability of iDKT parameters by executing a *cross-model simulation*. We "inject" the projected iDKT parameters (\mathbf{l}_u and \mathbf{t}_u) into the canonical BKT recurrence equations to generate *induced mastery trajectories*. This allows us to verify whether the learned factors preserve their causal roles. The metric is statistically quantified as the Pearson correlation coefficient between these induced trajectories and the theoretical reference trajectories generated by the original BKT model:

$$I_2 = \text{Corr}\left(P(L)_{\text{ind}}, P(L)_{\text{ref}}\right) \quad (15)$$

where $P(L)_{\text{ind}}$ and $P(L)_{\text{ref}}$ represent the sequences of mastery probabilities at each timestep for the induced and reference models, respectively.

Table 2 shows interpretability metrics for the unconstrained iDKT model ($\lambda_{\text{ref}} = 0$). Upon enabling Representational Grounding ($\lambda_{\text{ref}} = 0.10$), we observe a significant restoration of semantic alignment, as detailed in Table 3.

Table 2. Interpretability Alignment Metrics for unconstrained iDKT ($\lambda_{ref} = 0$).

Dataset	I_1 (Init.)	Int.	I_1 (Rate)	Int.	I_2	Int.
AS2009	-0.1382	Poor	-0.0067	Negl.	0.1870	Poor
AS2015	-0.0392	Negl.	0.0908	Negl.	0.0975	Negl.
Algebra2005	-0.0602	Negl.	-0.0273	Negl.	0.0600	Negl.
Bridge2006	-0.0645	Negl.	0.0160	Negl.	0.0809	Negl.
NIPS34	0.2070	Poor	-0.0425	Negl.	0.1722	Poor

Table 3. Interpretability Alignment Metrics for grounded iDKT ($\lambda_{ref} = 0.1$).

Dataset	I_1 (Init.)	Int.	I_1 (Rate)	Int.	I_2	Int.
AS2009_S	0.5409	Good	0.3131	Fair	0.6250	Excell.
AS2015	0.9217	Excell.	0.8801	Excell.	0.1749	Poor
Algebra2005	0.5444	Good	0.3310	Fair	0.0661	Negl.
Bridge2006	0.4561	Fair	0.6422	Good	0.0859	Negl.

The results presented in Tables 2 and 3 reveal a notable divergence between *Convergent Validity* (I_1) and *Predictor Equivalence* (I_2). Although I_1 attains high levels across most datasets even with relatively low λ_{ref} grounding weights, I_2 remains consistently low.

High I_1 scores demonstrate that iDKT successfully internalizes the *semantic identity* of the theoretical constructs (Initial Knowledge and Learning Rate). However, I_2 measures functional substitutability within a comparatively constrained reference model that is unable to capture the intricate behavioral patterns identified by iDKT. As a high-capacity Transformer, iDKT captures complex, long-range dependencies and context-aware dynamics that exceed the modeling capability of classical BKT. This discrepancy provides empirical evidence that iDKT does not merely mimic the reference model, but rather maps its core pedagogical constructs onto a more sophisticated and predictive architecture.

6.6. Accuracy–Interpretability Trade-Off (RQ2)

Prediction performance and interpretability are considered as two opposing goals. Our approach to explore such trade-off is based in building a Pareto frontier, which is constructed by gradually increasing the grounding weight λ_{ref} defined in Equation ?? and plotting the resulting values of performance and interpretability.

We measure interpretability using the *Composite Alignment I* metric, defined as the arithmetic mean of the individual alignment components I_1 and I_2 defined in Equations 13 and 15:

$$I = \frac{1}{3}(I_{1l} + I_{1t} + I_2) \quad (16)$$

Tables 4 and 6 summarize the performance and alignment results for the ASSISTments 2009 and 2015 datasets across a systematic sweep of the grounding weight λ_{ref} (Equation ??).

The resulting Pareto curves are visualized in Figures 5 and 7. We observe that theoretical guidance serves as a powerful regularizer; even a small weight ($\lambda = 0.10$) causes a massive leap in interpretability with minimal impact on accuracy. However, excessive grounding leads to a collapse in denser datasets like AS2009, where the model can no longer reconcile complex behavioral patterns with simple pedagogical assumptions.

6.7. Improvement of Pedagogical Diagnostics (RQ3)

In the following sections, we analyze how the architectural design of iDKT can be leveraged to improve pedagogical diagnostics

Table 4. Pareto Sweep results for ASSISTments 2015.

λ_{ref}	Performance (AUC)	Interpretability (I)	Interpretation
0.0 (Baseline)	0.7252	0.0497	Black Box
0.1 (Sweet Spot)	0.7233	0.6589	99.7% AUC Retained
0.2	0.7169	0.6709	Peak Alignment
0.3	0.7093	0.6639	Saturated
0.5	0.6942	0.6362	Saturated
1.0	0.6724	0.5519	Theory-Locked

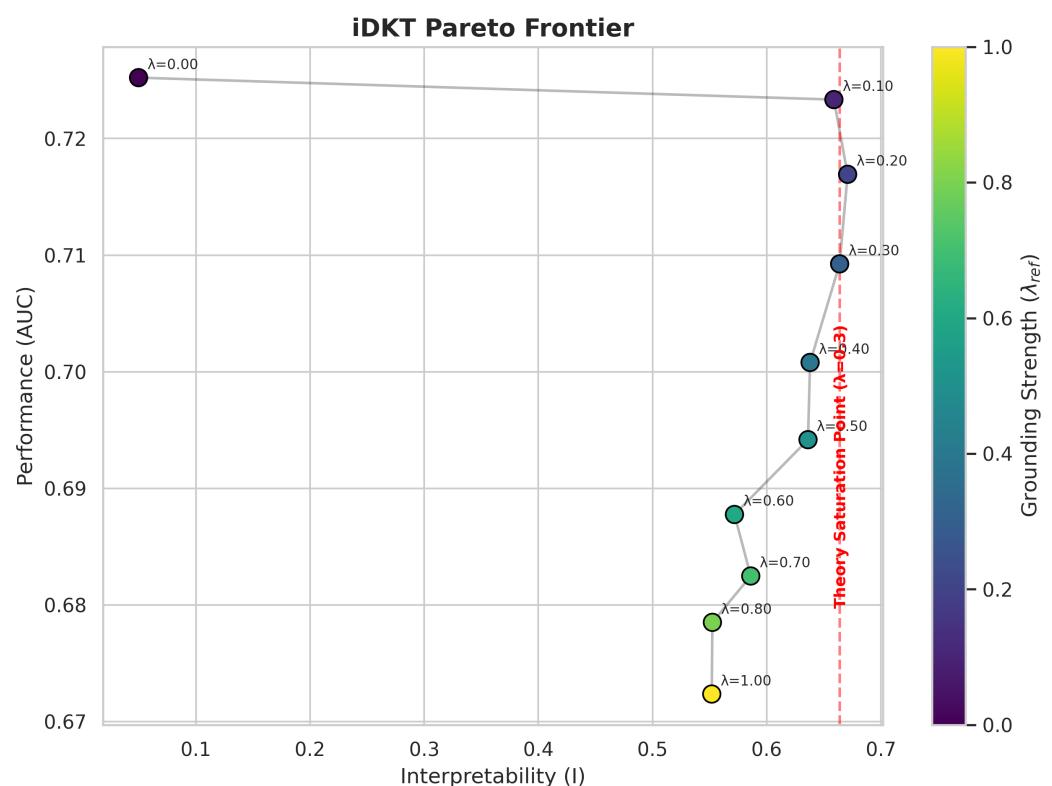
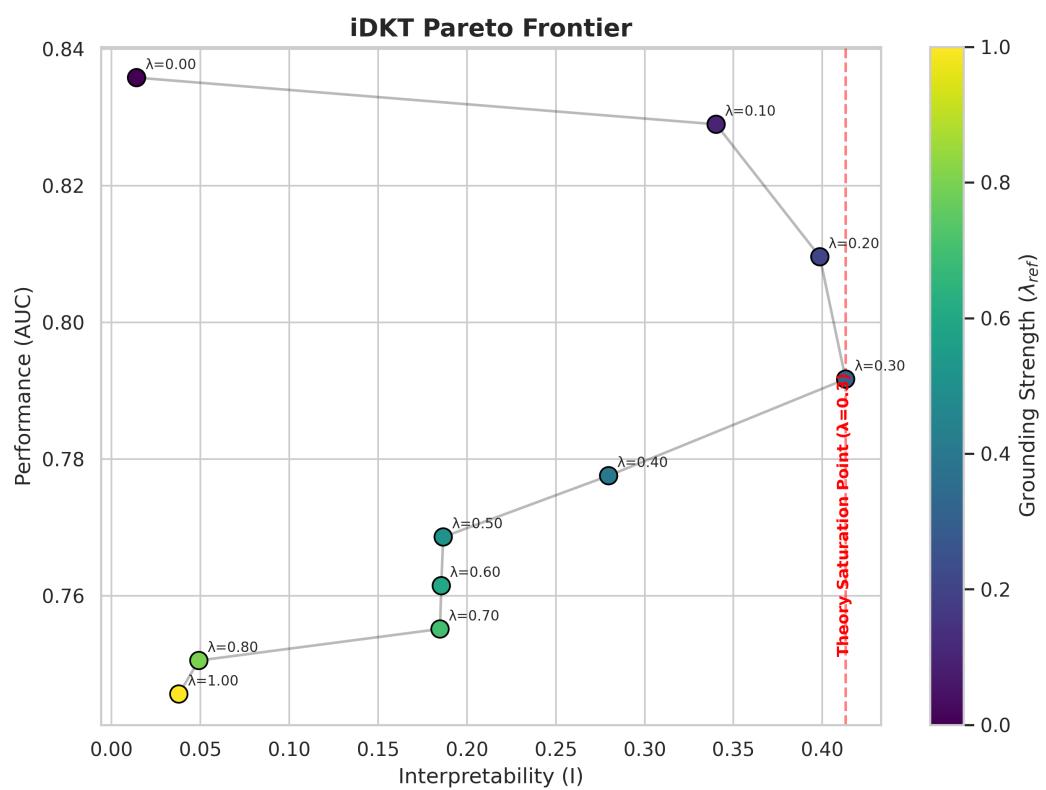
**Table 5.** iDKT Pareto Frontier for ASSISTments 2015. The trajectory illustrates a saturation of theoretical alignment at $\lambda \approx 0.2$, followed by a drift at higher weights where performance is traded without gaining interpretability.

Table 6. Pareto Sweep results for ASSISTments 2009.

λ_{ref}	Performance (AUC)	Interpretability (I)	Interpretation
0.0 (Baseline)	0.8358	0.0140	Black Box
0.1 (Sweet Spot)	0.8290	0.3402	Accurate & Balanced
0.2	0.8096	0.3986	Theory-Dominant
0.3	0.7917	0.4131	Peak Alignment
0.5	0.7686	0.1866	Saturated
1.0	0.7456	0.0378	Theoretical Collapse

**Table 7.** iDKT Pareto Frontier for ASSISTments 2009. The curve shows a non-monotonic relationship where interpretability peaks at $\lambda \approx 0.3$, followed by a *grounding collapse* as the model fails to reconcile excessive constraints.

6.7.1. Contextualization

The capability of iDKT to leverage longitudinal context is illustrated in Figure 2, which contrasts its mastery estimations with the BKT baseline for a high-proficiency learner. While BKT leads to excessive volatility—where isolated failures are interpreted as significant knowledge drops—iDKT maintains a stable and resilient diagnostic trajectory.

The core of this advantage lies in the self-attention mechanism, which enables iDKT to encode longitudinal context by dynamically weighting the entire interaction history. Instead of relying solely on the most recent state, iDKT evaluates each interaction in the context of the student's established behavioral patterns. In the scenario shown in Figure 2, the model maintains a high mastery estimation despite sporadic failures because the attention weights remain anchored to the student's longitudinal progression. This capacity to distinguish between temporary behavioral noise (e.g., slips) and genuine changes in knowledge state represents a fundamental improvement over classical baselines, supporting robust and individualized diagnostics.

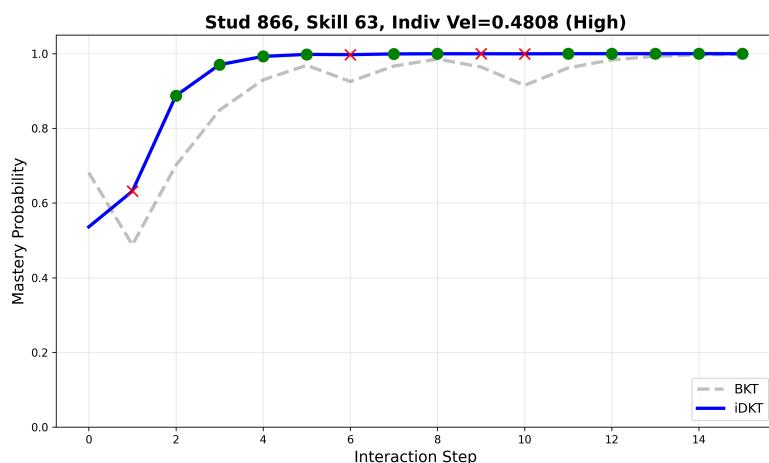


Figure 2. Comparison of two mastery trajectories estimated by iDKT and BKT illustrating how iDKT can be more stable and resilient against behavioral noise.

6.7.2. Individualization

The primary value of iDKT lies in its ability to transform population-level theoretical averages into high-granular contextual diagnostics. While standard BKT assumes that every student shares a fixed initial mastery (L_0) and learning rate (T) for a given skill, iDKT decomposes these into individualized profiles.

To statistically quantify this granularity, we analyze the *delta distribution* ($\Delta = t_s - T$), representing the student-specific deviation from the theoretical learning rate across all skills. In both ASSISTments datasets, we observe a significant non-zero standard deviation ($\sigma \approx 0.019$ for AS2009 and $\sigma \approx 0.039$ for AS2015). This represents a substantial increase in diagnostic resolution compared to the BKT baseline, where $\sigma = 0$. The right-skewed tail in the distributions identifies a sub-population of fast learners whose true acquisition pace is systematically underestimated by classical models.

Figure ?? visualizes the delta distributions for both datasets. The width of the curves quantifies the pedagogical information that is lost when using population-level averages. This variance allows the system to distinguish between students who are faster than suggested by theoretical priors and those who require additional practice to reach mastery.

6.7.3. High-Granularity Trajectory Analysis

To provide definitive visual evidence for RQ3, we generated a Mastery Alignment Mosaic (Figure 3) displaying individualized trajectories for three learning archetypes (Fast, Median, Slow) across 15 skills. For rigorous validation, we utilized the “Set S Isolation” strategy, restricting the analysis to students with continuous interaction histories (≤ 200 steps) and recalibrating the BKT reference exclusively on this population.

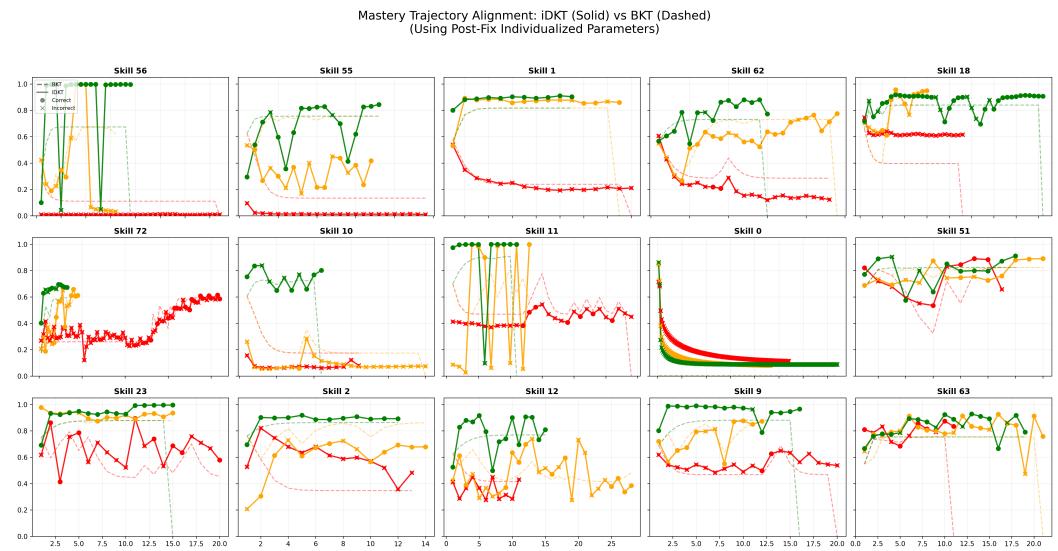


Figure 3. Mastery Trajectory Alignment Mosaic for AS2009_S. Solid lines represent iDKT individualized mastery, while dashed lines show the BKT baseline. Markers indicate observed student performance (Circle: Correct, X: Incorrect). Trajectories are selected from students with > 10 interactions per skill.

The qualitative analysis of these trajectories reveals three significant pedagogical insights:

1. **Dynamic Diagnostic Placement:** Even when BKT is forced to start at the population prior (L_0), iDKT starting points ($t = 1$) vary per student. This demonstrates that iDKT leverages cross-skill behavioral transfers to perform individualized placement before the first interaction with a new skill.
2. **Trajectory Crossovers:** We observe instances where students starting with lower initial mastery overtake peers due to a higher identified learning velocity (t_s). This confirms iDKT’s ability to model individualized acquisition rates beyond population-level curves.
3. **Diagnostic Resilience:** iDKT displays superior stability compared to BKT baselines. While isolated failures cause sharp drops in BKT dashed lines, iDKT solid lines for high-velocity learners remain stable, correctly identifying temporary slips vs. fundamental knowledge loss.

6.7.4. Confidence

Individualized student profiling offers new opportunities for pedagogical monitoring and adaptive intervention. A significant practical application is evaluating the reliability of predictions made by intrinsically interpretable models, such as BKT. By comparing BKT mastery estimations with those from iDKT, we can derive a measure of pedagogical confidence in the theoretical baseline.

This relationship is visualized in Figure 4, which displays the concordance between BKT and iDKT mastery predictions across the sub-population of the most frequent skills

and students for both datasets. The ASSISTments 2015 results demonstrate high levels of agreement, whereas the ASSISTments 2009 data reveals more divergence, where iDKT uncovers individualized patterns beyond the theoretical prior. Such visualizations provide actionable pedagogical insights that complement traditional baseline diagnostics. This visualization framework provides robust support for personalized longitudinal tracking by identifying students or curricular content that deviate from theoretical expectations. These divergences serve as critical diagnostic indicators for cases that may warrant prioritized pedagogical intervention.

392
393
394
395
396
397
398
399

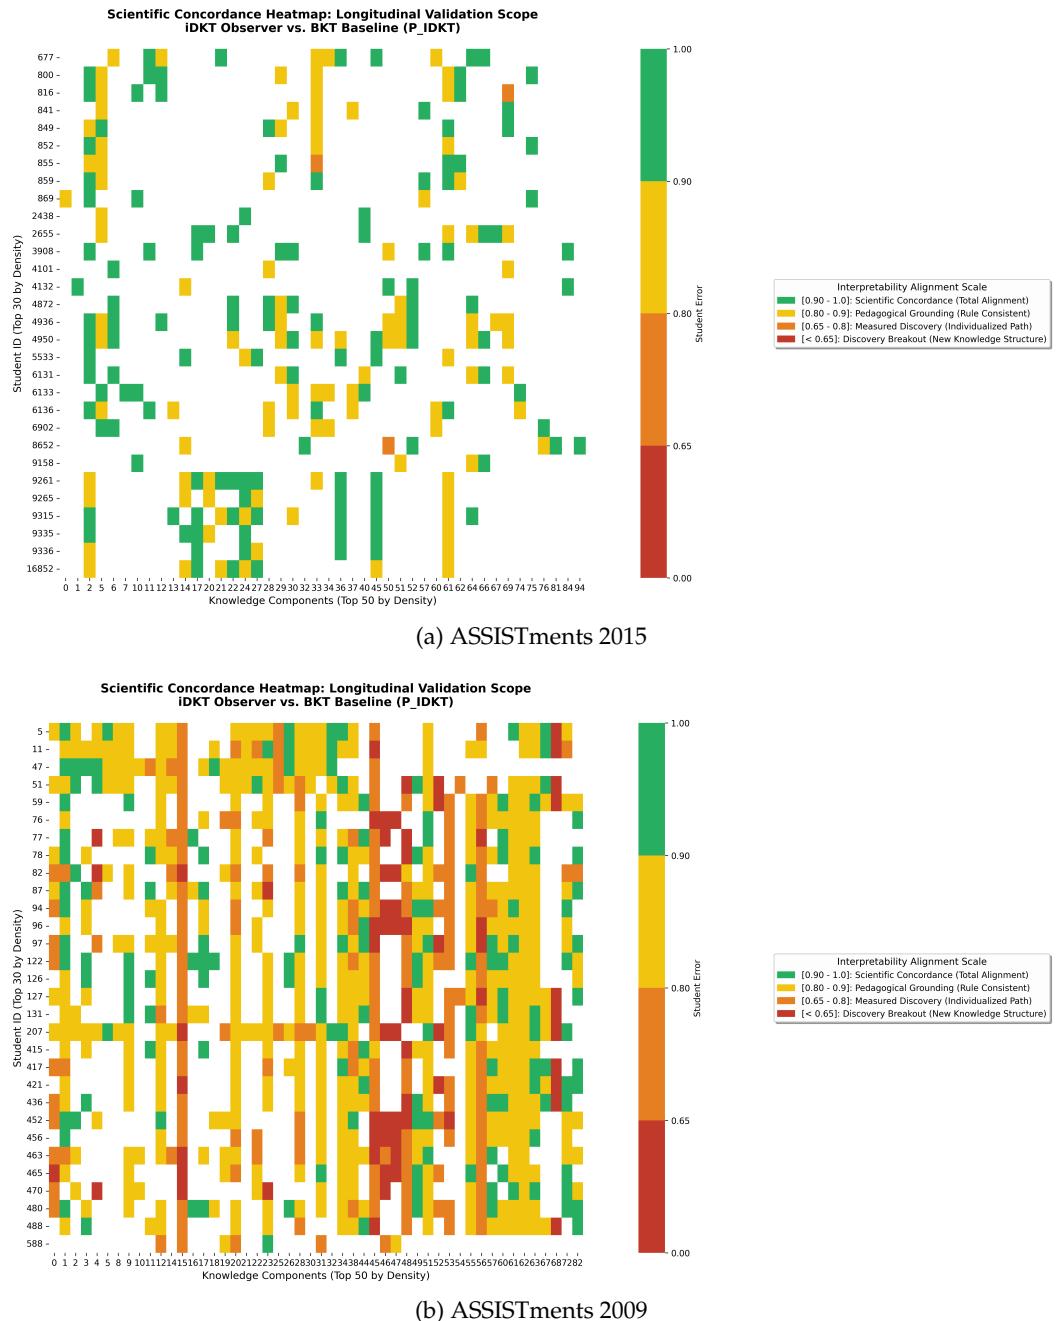


Figure 4. Heatmaps of per-skill prediction alignment.

7. Conclusions

This work has introduced iDKT, a Transformer-based Knowledge Tracing model that bridges the gap between the high predictive capacity of deep learning and the intrinsic in-

400
401
402

terpretability of models such as BKT, which utilize constructs with clear semantic meanings. By utilizing Representational Grounding, we have demonstrated that it is possible to anchor deep latent representations to semantically meaningful constructs without sacrificing state-of-the-art performance.

Our experimental results successfully address the research questions. First, we proved that iDKT successfully internalizes the theoretical constructs of the reference model, achieving high convergent validity (I_1). Second, we formalized a methodology for measuring the trade-off between predictive performance and interpretability, enabling the identification of the “Interpretability Sweet Spot” where significant alignment is achieved with minimal loss in predictive AUC. Finally, we demonstrated that iDKT provides a high increase in diagnostic granularity compared to population-level baselines by identifying individualized learning velocities that enable truly adaptive pacing.

The primary contribution of this research is twofold: a robust methodology for evaluating the internal interpretability of Transformer-based models in education, and a practical architecture that transforms “black-box” predictors into interpretable tools for knowledge tracing. By grounding deep learning in established educational concepts, iDKT offers a path toward AI-driven personalization that is both highly accurate and pedagogically actionable, providing a foundation for next-generation intelligent tutoring systems.

Author Contributions: Conceptualization, C. L.; methodology, C. L. and O. C. S.; software, C. L.; validation, C. L. and O. C. S.; formal analysis, C. L.; investigation, C. L.; resources, C. L. and O. C. S.; data curation, C. L.; writing—original draft preparation, C. L.; writing—review and editing, C. L. and O. C. S.; visualization, C. L.; supervision, O. C. S.; project administration, O. C. S.; funding acquisition, O. C. S. All authors have read and agreed to the published version of the manuscript.

Funding: This research received no external funding.

Data Availability Statement: The datasets used in this paper can be downloaded through the links provided in <https://pykt-toolkit.readthedocs.io/en/latest/datasets.html>.

Conflicts of Interest: The authors declare no conflicts of interest.

References

- Corbett, A.T.; Anderson, J.R. Knowledge tracing: Modeling the acquisition of procedural knowledge. *User modeling and user-adapted interaction* **1994**, *4*, 253–278.
- Šarić Grgić, I.; Grubišić, A.; Gašpar, A. Twenty-five years of Bayesian knowledge tracing: a systematic review, 2022.
- Piech, C.; Bassett, J.; Huang, J.; Ganguli, S.; Sahami, M.; Guibas, L.J.; Sohl-Dickstein, J. Deep knowledge tracing. *Advances in neural information processing systems* **2015**, *28*.
- Vaswani, A.; Shazeer, N.; Parmar, N.; Uszkoreit, J.; Jones, L.; Gomez, A.N.; Kaiser, Ł.; Polosukhin, I. Attention is all you need. *Advances in neural information processing systems* **2017**, *30*.
- Abdelrahman, G.; Wang, Q.; Nunes, B. Knowledge tracing: A survey. *ACM Computing Surveys* **2023**, *55*, 1–37.
- Bai, X.; et al. A Survey of Explainable Knowledge Tracing, 2024.
- Fantozzi, M.; et al. The Explainability of Transformers - Current Status and Directions. *arXiv preprint arXiv:2401.09202* **2024**.
- Di Marino, S.; et al. Ante-Hoc Methods for Interpretable Deep Models: A Survey, 2025.
- Karpatne, A.; Atluri, G.; Faghmous, J.; Steinbach, M.; Banerjee, A.; Ganguly, A.; Shekhar, S.; Samatova, N.; Kumar, V. Theory-guided Data Science: A New Paradigm for Scientific Discovery from Data. *IEEE Transactions on Knowledge and Data Engineering* **2017**, *29*, 2318–2331.
- Willard, J.; Jia, X.; Xu, S.; Steinbach, M.; Kumar, V. Integrating Physics-Based Modeling With Machine Learning: A Survey, 2022.
- Von Rueden, L.; Mayer, S.; Beckh, K.; Georgiev, B.; Giesselbach, S.; Heese, R.; Kirsch, B.; Pfrommer, J.; Pick, A.; Ramamurthy, R.; et al. Informed Machine Learning – A Taxonomy and Survey of Integrating Knowledge into Learning Systems. *IEEE Transactions on Knowledge and Data Engineering* **2021**, p. 1–1.
- Liu, Z.; Liu, Q.; Chen, J.; Huang, S.; Tang, J.; Luo, W. pyKT: a python library to benchmark deep learning based knowledge tracing models. *Advances in Neural Information Processing Systems* **2022**, *35*, 18542–18555.
- Feng, M.; Heffernan, N.; Koedinger, K. Addressing the assessment challenge with an online system that tutors as it assesses. *User modeling and user-adapted interaction* **2009**, *19*, 243–266.

-
- 14. Stamper, J.; Niculescu-Mizil, A.; Ritter, S.; Gordon, G.; Koedinger, K. Algebra I 2005-2006. *Development data sets from KDD Cup 2010.* 453
454
 - 15. Stamper, J.; Niculescu-Mizil, A.; Ritter, S.; Gordon, G.; Koedinger, K. Bridge to algebra 2006-2007. *Development data sets from KDD Cup 2010.* 455
456
 - 16. Wang, Z.; Lamb, A.; Saveliev, E.; Cameron, P.; Zaykov, Y.; Hernández-Lobato, J.M.; Turner, R.E.; Baraniuk, R.G.; Barton, C.; Jones, S.P.; et al. Instructions and guide for diagnostic questions: The neurips 2020 education challenge. *arXiv preprint arXiv:2007.12061 2020.* 457
458
459
 - 17. Campbell, D.T.; Fiske, D.W. Convergent and discriminant validation by the multitrait-multimethod matrix. *Psychological bulletin 1959, 56, 81.* 460
461
 - 18. Association, A.E.R.; Association, A.P.; on Measurement in Education, N.C.; et al. Standards for educational and psychological testing. (*No Title*) 1985. 462
463

Disclaimer/Publisher's Note: The statements, opinions and data contained in all publications are solely those of the individual author(s) and contributor(s) and not of MDPI and/or the editor(s). MDPI and/or the editor(s) disclaim responsibility for any injury to people or property resulting from any ideas, methods, instructions or products referred to in the content. 464
465
466